

# SESSION 10 - Project on API based Image Classification

1. Image classification
  2. Setup
  3. Download and explore the dataset
  4. Load data using a Keras utility
  5. Visualize the data
  6. Configure the dataset for performance
  7. Standardize the data
  8. A basic Keras model
    - 8.1 Create the model
    - 8.2 Compile the model
    - 8.3 Model summary
    - 8.4 Train the model
  9. Visualize training results
  10. Overfitting
  11. Data augmentation
  12. Dropout
  13. Compile and train the model
  14. Visualize training results
  15. Predict on new data
- 

## 1. Image classification

This tutorial shows how to classify images of flowers using a `tf.keras.Sequential` model and load data using `tf.keras.utils.image_dataset_from_directory`. It demonstrates the following concepts:

- Efficiently loading a dataset off disk.
- Identifying overfitting and applying techniques to mitigate it, including data augmentation and dropout.

This tutorial follows a basic machine learning workflow:

1. Examine and understand data
2. Build an input pipeline
3. Build the model
4. Train the model
5. Test the model
6. Improve the model and repeat the process

## 2. Setup

Import TensorFlow and other necessary libraries:

```
In [ ]: import matplotlib.pyplot as plt
import numpy as np
import PIL
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
```

## 3. Download and explore the dataset

This tutorial uses a dataset of about 3,700 photos of flowers. The dataset contains five sub-directories, one per class:

```
flower_photo/
  daisy/
  dandelion/
  roses/
  sunflowers/
  tulips/
```

```
In [ ]: import pathlib

dataset_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz"
```

```
data_dir = tf.keras.utils.get_file('flower_photos.tar', origin=dataset_url, extract=True)  
data_dir = pathlib.Path(data_dir).with_suffix('')
```

After downloading, you should now have a copy of the dataset available. There are 3,670 total images:

```
In [ ]: image_count = len(list(data_dir.glob('*/*.jpg')))  
        print(image_count)
```

3670

Here are some roses:

```
In [ ]: roses = list(data_dir.glob('roses/*'))  
        PIL.Image.open(str(roses[0]))
```

Out[ ]:



```
In [ ]: PIL.Image.open(str(roses[1]))
```

Out[ ]:



And some tulips:

In [ ]:

```
tulips = list(data_dir.glob('tulips/*'))  
PIL.Image.open(str(tulips[0]))
```

Out[ ]:



In [ ]:

```
PIL.Image.open(str(tulips[1]))
```

Out[ ]:



## 4. Load data using a Keras utility

Next, load these images off disk using the helpful `tf.keras.utils.image_dataset_from_directory` utility. This will take you from a directory of images on disk to a `tf.data.Dataset` in just a couple lines of code. If you like, you can also write your own data loading code from scratch by visiting the [Load and preprocess images](#) tutorial.

### Create a dataset

Define some parameters for the loader:

```
In [ ]: batch_size = 32
        img_height = 180
        img_width = 180
```

It's good practice to use a validation split when developing your model. Use 80% of the images for training and 20% for validation.

```
In [ ]: train_ds = tf.keras.utils.image_dataset_from_directory(
        data_dir,
```

```
validation_split=0.2,  
subset="training",  
seed=123,  
image_size=(img_height, img_width),  
batch_size=batch_size)
```

Found 3670 files belonging to 5 classes.  
Using 2936 files for training.

```
In [ ]: val_ds = tf.keras.utils.image_dataset_from_directory(  
    data_dir,  
    validation_split=0.2,  
    subset="validation",  
    seed=123,  
    image_size=(img_height, img_width),  
    batch_size=batch_size)
```

Found 3670 files belonging to 5 classes.  
Using 734 files for validation.

You can find the class names in the `class_names` attribute on these datasets. These correspond to the directory names in alphabetical order.

```
In [ ]: class_names = train_ds.class_names  
print(class_names)
```

```
['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips']
```

## 5. Visualize the data

Here are the first nine images from the training dataset:

```
In [ ]: import matplotlib.pyplot as plt  
  
plt.figure(figsize=(10, 10))  
for images, labels in train_ds.take(1):  
    for i in range(9):  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(images[i].numpy().astype("uint8"))  
        plt.title(class_names[labels[i]])  
        plt.axis("off")
```



roses



dandelion



tulips



sunflowers




dandelion



roses




dandelion



roses



tulips



You will pass these datasets to the Keras `Model.fit` method for training later in this tutorial. If you like, you can also manually iterate over the dataset and retrieve batches of images:

```
In [ ]: for image_batch, labels_batch in train_ds:
        print(image_batch.shape)
```

```
print(labels_batch.shape)
break
```

```
(32, 180, 180, 3)
(32,)
```

The `image_batch` is a tensor of the shape `(32, 180, 180, 3)`. This is a batch of 32 images of shape `180x180x3` (the last dimension refers to color channels RGB). The `label_batch` is a tensor of the shape `(32,)`, these are corresponding labels to the 32 images.

You can call `.numpy()` on the `image_batch` and `labels_batch` tensors to convert them to a `numpy.ndarray`.

## 6. Configure the dataset for performance

Make sure to use buffered prefetching, so you can yield data from disk without having I/O become blocking. These are two important methods you should use when loading data:

- `Dataset.cache` keeps the images in memory after they're loaded off disk during the first epoch. This will ensure the dataset does not become a bottleneck while training your model. If your dataset is too large to fit into memory, you can also use this method to create a performant on-disk cache.
- `Dataset.prefetch` overlaps data preprocessing and model execution while training.

Interested readers can learn more about both methods, as well as how to cache data to disk in the *Prefetching* section of the [Better performance with the tf.data API](#) guide.

In [ ]:

```
AUTOTUNE = tf.data.AUTOTUNE

train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

## 7. Standardize the data

The RGB channel values are in the `[0, 255]` range. This is not ideal for a neural network; in general you should seek to make your input values small.

Here, you will standardize values to be in the `[0, 1]` range by using `tf.keras.layers.Rescaling`:

In [ ]:

```
normalization_layer = layers.Rescaling(1./255)
```



There are two ways to use this layer. You can apply it to the dataset by calling `Dataset.map` :

```
In [ ]: normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
image_batch, labels_batch = next(iter(normalized_ds))
first_image = image_batch[0]
# Notice the pixel values are now in `[0,1]`.
print(np.min(first_image), np.max(first_image))
```

```
0.0 1.0
```

Or, you can include the layer inside your model definition, which can simplify deployment. Use the second approach here.

Note: You previously resized images using the `image_size` argument of `tf.keras.utils.image_dataset_from_directory` . If you want to include the resizing logic in your model as well, you can use the `tf.keras.layers.Resizing` layer.

## 8. A basic Keras model

### 8.1 Create the model

The Keras `Sequential` model consists of three convolution blocks ( `tf.keras.layers.Conv2D` ) with a max pooling layer ( `tf.keras.layers.MaxPooling2D` ) in each of them. There's a fully-connected layer ( `tf.keras.layers.Dense` ) with 128 units on top of it that is activated by a ReLU activation function ( `'relu'` ). This model has not been tuned for high accuracy; the goal of this tutorial is to show a standard approach.

```
In [ ]: num_classes = len(class_names)

model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])
```

### 8.2 Compile the model

For this tutorial, choose the `tf.keras.optimizers.Adam` optimizer and `tf.keras.losses.SparseCategoricalCrossentropy` loss function. To view training and validation accuracy for each training epoch, pass the `metrics` argument to `Model.compile`.

```
In [ ]: model.compile(optimizer='adam',
                    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                    metrics=['accuracy'])
```

### 8.3 Model summary

View all the layers of the network using the Keras `Model.summary` method:

```
In [ ]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling_1 (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d (MaxPooling2D)	(None, 90, 90, 16)	0
conv2d_1 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 45, 45, 32)	0
conv2d_2 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 22, 22, 64)	0
flatten (Flatten)	(None, 30976)	0
dense (Dense)	(None, 128)	3965056
dense_1 (Dense)	(None, 5)	645
Total params: 3989285 (15.22 MB)		
Trainable params: 3989285 (15.22 MB)		

Non-trainable params: 0 (0.00 Byte)

---

## 8.4 Train the model

Train the model for 10 epochs with the Keras `Model.fit` method:

In [ ]:

```
epochs=10
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)
```

Epoch 1/10

92/92 [=====] - 64s 672ms/step - loss: 1.2298 - accuracy: 0.4809 - val\_loss: 1.0611 - val\_accuracy: 0.5640

Epoch 2/10

92/92 [=====] - 58s 621ms/step - loss: 0.9407 - accuracy: 0.6304 - val\_loss: 0.9395 - val\_accuracy: 0.6403

Epoch 3/10

92/92 [=====] - 52s 566ms/step - loss: 0.7508 - accuracy: 0.7115 - val\_loss: 0.8593 - val\_accuracy: 0.6608

Epoch 4/10

92/92 [=====] - 49s 532ms/step - loss: 0.4982 - accuracy: 0.8174 - val\_loss: 0.9871 - val\_accuracy: 0.6608

Epoch 5/10

92/92 [=====] - 44s 480ms/step - loss: 0.2844 - accuracy: 0.9022 - val\_loss: 1.0136 - val\_accuracy: 0.6717

Epoch 6/10

92/92 [=====] - 51s 557ms/step - loss: 0.1572 - accuracy: 0.9533 - val\_loss: 1.2876 - val\_accuracy: 0.6553

Epoch 7/10

92/92 [=====] - 57s 615ms/step - loss: 0.0988 - accuracy: 0.9721 - val\_loss: 1.4480 - val\_accuracy: 0.6757

Epoch 8/10

92/92 [=====] - 61s 657ms/step - loss: 0.0484 - accuracy: 0.9881 - val\_loss: 1.7409 - val\_accuracy: 0.6417

Epoch 9/10

92/92 [=====] - 54s 580ms/step - loss: 0.0351 - accuracy: 0.9908 - val\_loss: 1.7087 - val\_accuracy: 0.6676

Epoch 10/10

92/92 [=====] - 53s 571ms/step - loss: 0.0386 - accuracy: 0.9881 - val\_loss: 1.8021 - val\_accuracy: 0.6703

## 9. Visualize training results

Create plots of the loss and accuracy on the training and validation sets:

```
In [ ]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



The plots show that training accuracy and validation accuracy are off by large margins, and the model has achieved only around 60% accuracy on the validation set.

The following tutorial sections show how to inspect what went wrong and try to increase the overall performance of the model.

## 10. Overfitting

In the plots above, the training accuracy is increasing linearly over time, whereas validation accuracy stalls around 60% in the training process. Also, the difference in accuracy between training and validation accuracy is noticeable—a sign of [overfitting](#).

When there are a small number of training examples, the model sometimes learns from noises or unwanted details from training examples—to an extent that it negatively impacts the performance of the model on new examples. This phenomenon is known as overfitting. It means that the model will have a difficult time generalizing on a new dataset.

There are multiple ways to fight overfitting in the training process. In this tutorial, you'll use *data augmentation* and add *dropout* to your model.

## 11. Data augmentation

Overfitting generally occurs when there are a small number of training examples. [Data augmentation](#) takes the approach of generating additional training data from your existing examples by augmenting them using random transformations that yield believable-looking images. This helps expose the model to more aspects of the data and generalize better.

You will implement data augmentation using the following Keras preprocessing layers: `tf.keras.layers.RandomFlip`, `tf.keras.layers.RandomRotation`, and `tf.keras.layers.RandomZoom`. These can be included inside your model like other layers, and run on the GPU.

```
In [ ]: data_augmentation = keras.Sequential(  
    [  
        layers.RandomFlip("horizontal",  
                           input_shape=(img_height,  
                                         img_width,  
                                         3)),  
        layers.RandomRotation(0.1),  
        layers.RandomZoom(0.1),  
    ]  
)
```

Visualize a few augmented examples by applying data augmentation to the same image several times:

```
In [ ]: plt.figure(figsize=(10, 10))  
for images, _ in train_ds.take(1):  
    for i in range(9):  
        augmented_images = data_augmentation(images)  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(augmented_images[0].numpy().astype("uint8"))  
        plt.axis("off")
```





You will add data augmentation to your model before training in the next step.

## 12. Dropout

Another technique to reduce overfitting is to introduce [dropout{:.external}](#) regularization to the network.

When you apply dropout to a layer, it randomly drops out (by setting the activation to zero) a number of output units from the layer during the training process. Dropout takes a fractional number as its input value, in the form such as 0.1, 0.2, 0.4, etc. This means dropping out 10%, 20% or 40% of the output units randomly from the applied layer.

Create a new neural network with `tf.keras.layers.Dropout` before training it using the augmented images:

```
In [ ]: model = Sequential([
    data_augmentation,
    layers.Rescaling(1./255),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes, name="outputs")
])
```

### 13. Compile and train the model

```
In [ ]: model.compile(optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'])
```

```
In [ ]: model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
sequential_1 (Sequential)	(None, 180, 180, 3)	0
rescaling_2 (Rescaling)	(None, 180, 180, 3)	0
conv2d_3 (Conv2D)	(None, 180, 180, 16)	448

max_pooling2d_3 (MaxPoolin g2D)	(None, 90, 90, 16)	0
conv2d_4 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_4 (MaxPoolin g2D)	(None, 45, 45, 32)	0
conv2d_5 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_5 (MaxPoolin g2D)	(None, 22, 22, 64)	0
dropout (Dropout)	(None, 22, 22, 64)	0
flatten_1 (Flatten)	(None, 30976)	0
dense_2 (Dense)	(None, 128)	3965056
outputs (Dense)	(None, 5)	645

```
=====
Total params: 3989285 (15.22 MB)
Trainable params: 3989285 (15.22 MB)
Non-trainable params: 0 (0.00 Byte)
```

In [ ]:

```
epochs = 15
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)
```

Epoch 1/15

```
92/92 [=====] - 59s 631ms/step - loss: 1.3787 - accuracy: 0.4016 - val_loss: 1.0822 - val_accuracy: 0.5831
```

Epoch 2/15

```
92/92 [=====] - 61s 668ms/step - loss: 1.0573 - accuracy: 0.5698 - val_loss: 1.0560 - val_accuracy: 0.5736
```

Epoch 3/15

```
92/92 [=====] - 63s 675ms/step - loss: 0.9716 - accuracy: 0.6144 - val_loss: 0.9455 - val_accuracy: 0.6063
```

Epoch 4/15

```
92/92 [=====] - 55s 602ms/step - loss: 0.8860 - accuracy: 0.6570 - val_loss: 0.8648 - val_accuracy: 0.6471
```

```

Epoch 5/15
92/92 [=====] - 63s 685ms/step - loss: 0.8106 - accuracy: 0.6897 - val_loss: 0.8483 - val_accuracy: 0.6649
Epoch 6/15
92/92 [=====] - 64s 695ms/step - loss: 0.7732 - accuracy: 0.7044 - val_loss: 0.8060 - val_accuracy: 0.6975
Epoch 7/15
92/92 [=====] - 62s 668ms/step - loss: 0.7373 - accuracy: 0.7183 - val_loss: 0.7856 - val_accuracy: 0.6853
Epoch 8/15
92/92 [=====] - 58s 631ms/step - loss: 0.7182 - accuracy: 0.7275 - val_loss: 0.8539 - val_accuracy: 0.6649
Epoch 9/15
92/92 [=====] - 92s 1s/step - loss: 0.6744 - accuracy: 0.7432 - val_loss: 0.7482 - val_accuracy: 0.7016
Epoch 10/15
92/92 [=====] - 82s 888ms/step - loss: 0.6465 - accuracy: 0.7514 - val_loss: 0.7273 - val_accuracy: 0.7125
Epoch 11/15
92/92 [=====] - 63s 682ms/step - loss: 0.6302 - accuracy: 0.7718 - val_loss: 0.7297 - val_accuracy: 0.7071
Epoch 12/15
92/92 [=====] - 65s 711ms/step - loss: 0.6053 - accuracy: 0.7708 - val_loss: 0.7093 - val_accuracy: 0.7289
Epoch 13/15
92/92 [=====] - 71s 772ms/step - loss: 0.5880 - accuracy: 0.7803 - val_loss: 0.7259 - val_accuracy: 0.7112
Epoch 14/15
92/92 [=====] - 78s 845ms/step - loss: 0.5469 - accuracy: 0.7933 - val_loss: 0.7747 - val_accuracy: 0.6935
Epoch 15/15
92/92 [=====] - 88s 962ms/step - loss: 0.5258 - accuracy: 0.8035 - val_loss: 0.7736 - val_accuracy: 0.7139

```

## 14. Visualize training results

After applying data augmentation and `tf.keras.layers.Dropout`, there is less overfitting than before, and training and validation accuracy are closer aligned:

```

In [ ]: acc = history.history['accuracy']
        val_acc = history.history['val_accuracy']


        loss = history.history['loss']
        val_loss = history.history['val_loss']

```

```
epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



## 15. Predict on new data

Use your model to classify an image that wasn't included in the training or validation sets.

Note: Data augmentation and dropout layers are inactive at inference time.

```
In [ ]: sunflower_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/592px-Red_sunflower.jpg"
sunflower_path = tf.keras.utils.get_file('Red_sunflower', origin=sunflower_url)

img = tf.keras.utils.load_img(
    sunflower_path, target_size=(img_height, img_width)
)
```



```
img_array = tf.keras.utils.img_to_array(img)
img_array = tf.expand_dims(img_array, 0) # Create a batch

predictions = model.predict(img_array)
score = tf.nn.softmax(predictions[0])

print(
    "This image most likely belongs to {} with a {:.2f} percent confidence."
    .format(class_names[np.argmax(score)], 100 * np.max(score))
)
```

Downloading data from [https://storage.googleapis.com/download.tensorflow.org/example\\_images/592px-Red\\_sunflower.jpg](https://storage.googleapis.com/download.tensorflow.org/example_images/592px-Red_sunflower.jpg)  
117948/117948 [=====] - 0s 2us/step  
1/1 [=====] - 2s 2s/step  
This image most likely belongs to sunflowers with a 99.96 percent confidence.

---